

- Dynamic Programming
- Wealth and Consumption Choice – A Cake Eating Problem
- Essential Reading:
 - Adda and Cooper Dynamic Economics: Quantitative Methods and Applications - Chapter 1. (See Moodle PDF)
 - Gregory, Chow Dynamic Economics: Optimization by the Lagrange Method - Chapter 2

Dynamic Programming

Dynamic Programming - Introduction

- In the last lecture, we used **Lagrange multipliers** to solve the **optimisation** problem of the firm.
- **Dynamic Programming** is an alternate method that can be used to solve optimisation problems.
- Developed in the 1940s by **Richard Bellman** at RAND Corporation
- Solves **multistage decision-making problems** by decomposing into smaller subproblems
- The approach is different yet gives an **identical solution**
- The name was chosen to avoid words like “**research, planning**” but still related to decision making, hence the credibility of “**Dynamic**” from Physics with the somewhat uninformative “**Programming**”
- something like **Chained Decomposition Solution Method** might make more sense

Dynamic Programming - Introduction


- **Dynamic Programming** is popular since it is **easy to implement numerically** with a computer.
- Widely used across Econ, Finance, Computer Science, Operations, Engineering, Game Theory, Machine Learning ... basically any quantitative field
- Most **modern macro(-finance) models** written in **recursive** (self-similar) form, essential tool
- We will learn dynamic programming using an **example: Eating a Cake**

Eating a Cake

Dynamic Programming with Cake-Eating Example

- Suppose you have a cake of size W_1 , **your wealth**. You have T **periods to consume** this cake.
- Every period $t = 1, 2, \dots, T$ you **consume some of the cake** and **save the rest**. The initial size of the cake at $t = 1$ is W_1 .
- Assume that the cake **cannot melt (depreciate) or grow**.
- Let c_t represent the consumption of cake at time t and $u(c_t)$ the flow of utility (satisfaction) from this consumption.
- Assume $u(\cdot)$ is real-valued, **continuous, differentiable and concave** and consumption should always be non-negative.
- Examples: $c^{0.5}$, $\ln(c)$, $\frac{c^{1-\sigma}}{1-\sigma}$. Might need to be careful with negative u in context

Example of Utility Function



- increasing and concave (the next piece is not as good as the last)

- The **life-time utility** from consuming the cake is given by the **discounted sum** of all current and future utility of consumption:

$$u(c_1) + \beta u(c_2) + \beta^2 u(c_3) + \dots + \beta^{T-1} u(c_T) \quad (1)$$

- That is,

$$\sum_{t=1}^T \beta^{t-1} u(c_t) \quad (2)$$

where $0 \leq \beta \leq 1$ is the **discount factor**, a measure of (im)patience.

Law of Motion of Cake

- The **evolution of cake** size (a.k.a law of motion) every period is given by::

$$W_{t+1} = W_t - c_t \quad (3)$$

Problem: How would you find the optimal path of consumption $\{c_t\}_{t=1}^T$

- In other words, what is the level of **consumption every period** that **maximizes your lifetime utility** in equation (2) above.
- We are looking for a consumption **plan** for all periods jointly: $\{c_t\}_{t=1}^T$.

Cake-Eating Example - Sequential Lagrangian Approach

- One approach is to use the method of Lagrange multipliers.
- This is then a constrained optimization problem where:

$$\max_{\{c_t, W_{t+1}\}_{t=1}^T} \left[\sum_{t=1}^T \beta^{t-1} u(c_t) \right]$$

- subject to the constraint:

$$W_{t+1} = W_t - c_t$$

for all $t = 1, 2, \dots T$.

Cake-Eating Example - Sequential Lagrangian Approach

The Lagrangian function can be written as:

$$\mathcal{L} = \sum_{t=1}^T \beta^{t-1} [u(c_t) - \lambda_t(W_{t+1} - W_t + c_t)]$$

Note: This is a dynamic optimization problem, we have an objective function and a constraint at every period t . All future values need to be discounted.

$$\begin{aligned} \mathcal{L} = & [u(c_1) + \beta u(c_2) + \dots + \beta^{t-1} u(c_t) + \dots \\ & - \lambda_1(W_2 - W_1 + c_1) - \beta \lambda_2(W_3 - W_2 + c_2) \dots \\ & - \beta^{t-1} \lambda_t(W_{t+1} - W_t + c_t) - \beta^t \lambda_{t+1}(W_{t+2} - W_{t+1} + c_{t+1}) \dots \end{aligned} \quad (4)$$

Remember! Like the Tobin model, we have to check for $t + 1$ -variables in two places

Cake-Eating Example - Sequential Lagrangian Approach

The necessary condition for maximizing this lagrangian function is given by the three FOCs:

$$\frac{\partial L}{\partial c_t} = 0 \Rightarrow u'(c_t) = \lambda_t$$

$$\frac{\partial L}{\partial W_{t+1}} = 0 \Rightarrow \lambda_t = \beta \lambda_{t+1}$$

$$\frac{\partial L}{\partial \lambda_t} = 0 \Rightarrow W_{t+1} = W_t - c_t$$

Cake-Eating Example - Euler Equation Intuition

- From eqs (1), (2), we get the **Euler equation**, intertemporal optimality condition:

$$\boxed{u'(c_t) = \beta u'(c_{t+1})} \quad (\text{EE})$$

- LHS represents the **marginal loss in utility** when you sacrifice a **small unit of consumption** and the RHS is the **discounted marginal gain in utility** from this extra unit of consumption next period.
- If the Euler equation holds, then it is **impossible to increase utility** by moving consumption across adjacent periods given a candidate solution $\{\tilde{c}_t\}_1^T = \{c_t^*\}_1^T$.
- No Arbitrage condition:** $+\beta u_{c,t+1}dc - u_{c,t}dc = 0$ or $-\beta u_{c,t+1}dc + u_{c,t}dc = 0$ depending on which way transfer consumption (small dc)

Euler Equation links periods $(t, t + 1)$

- From eqs (1), (2), we get the **Euler equation, intertemporal optimality condition**, for $t = 1, \dots, T - 1$:

$$u'(c_t) = \beta u'(c_{t+1})$$

- Example: if $u(c) = \ln(c)$ implies a (negative) growth path:

$$\frac{1}{c_t} = \beta \frac{1}{c_{t+1}} \Rightarrow \frac{c_{t+1} - c_t}{c_t} = (\% \text{ Growth in } c) = -(1 - \beta)$$

- Euler Equation makes a chain of pairs:


$$(c_1, c_2), (c_2, c_3), (c_3, c_4), \dots, (c_{99}, c_{100}) \quad (5)$$

$$(c_1, \beta c_1), (\beta c_1, \beta^2 c_1), (\beta^2 c_1, \beta^3 c_1), \dots, (\beta^{98} c_1, \beta^{99} c_1) \quad (6)$$

$$c_t = \beta^{t-1} c_1 \quad (7)$$

But we don't yet know c_1 . Once we solve for that, we get the full chain

Euler Equation links periods $(t, t + 1)$



Cake-Eating Example - Sequential Lagrangian Approach

Since this is a **finite time horizon** problem, we need to have a **terminal condition**.

For **maximum utility**, there should not be any cake left over at the end of the last period (no waste). That is,

$$W_{T+1} = 0 \quad (\text{END})$$

This terminal condition naturally implies that the sum of consumption across all periods should equal the total size of the cake (resource constraint, RC):

$$\sum_{t=1}^T c_t = W_1 \text{ (e.g. = 100)} \quad (\text{RC})$$

Using the value of W_1 (RC) and eq.s (EE) and (END), we can find the optimal path of consumption $\{c_t^*\}_{t=1}^T$ that maximizes utility.

For log-utility we can use pen and paper

We can plug the Euler Equation bridges into consumption, and use RC:

$$\sum_{t=1}^T c_t = W_1 \Rightarrow \sum_{t=1}^T \beta^{t-1} c_1 = W_1$$

We can arrange the sum:

$$c_1(1 + \beta + \beta^2 + \dots + \beta^{T-1}) = W_1$$

This is a geometric sum, we know from the toolkit how to solve this:

$$c_1 \frac{(1 - \beta^T)}{1 - \beta} = W_1$$

Solving for c as a function of parameters for patience and total periods:

$$c_1 = \frac{(1 - \beta)}{1 - \beta^T} W_1$$

And this **nects** the well-known **infinite horizon solution** ($T \rightarrow \infty$)

$$c_t = (1 - \beta)W_t \quad \forall t$$

Consume (e.g.) 5 percent of remaining cake (like every period is like the start)

For all the rest: Shooting Algorithm

1. **Initial Condition:** Start with W_1 , e.g. 100.
2. **Update:** Use Euler Equation in terms of cake, **second-order difference eqn:**

$$u_c(W_t - W_{t+1}) = \beta u_c(W_{t+1} - W_{t+2})$$

3. Rearrange, and guess W_2 :

$$W_{t+2} = W_{t+1} - u_c^{-1}((1/\beta)u_c(W_t - W_{t+1}))$$

4. Start: We have W_1 , guess W_2 , this implies W_3 . Then we can roll forward to get W_4, \dots, W_{T+1} . This is the first shot. Aim for zero.
5. **Terminal condition** Adjust guess W_2 , keep shooting until $W_{T+1} \approx 0$.
6. **Optimal Consumption** path: $C_t = W_{t-1} - W_t$
 - Fast numerical methods in **matlab**, **julia** etc to solve (Bisection!)
 - One can also do a **reverse shot**: We know $W_{T+1} = 0$, guess W_T to imply W_{T-1}, \dots, W_1 , and aim for starting $W_1 = 100$.

Time for Some Drawing!

Consumption and Wealth Sequences

Blue line: impatient; Purple: patient




Figure 1: $\{W_{t+1}, c_t\}_{t=1}^{T=20}$, $W_1 = 100$, $u(c) = \frac{c^{1-\sigma}}{1-\sigma}$

$$W_{t+2} = (1 + \beta^{1/\sigma})W_{t+1} - (\beta^{1/\sigma})W_t; \text{ given } W_{T+1} = 0; W_1 = 100$$

Cake-Eating Example - Value Function

- The solution to this T -period cake eating problem is found by substituting the optimal path of consumption in the lifetime utility function.
- We will denote this maximum as $V^T(W_1)$:

$$V^T(W_1) = \max \left[\sum_{t=1}^T \beta^{t-1} u(c_t) \right] = \sum_{t=1}^T \beta^{t-1} u(c_t^*)$$

- $V(W_1)$ is called as a **value function** and here it represents the **maximum T period utility of consumption** given an initial level of cake size W_1 .

Cake-Eating Example - Dynamic Programming Approach

- Suppose we change this cake eating problem by **adding a period 0** and giving an initial cake size of W_0 .
- We can again solve this by formulating a **new Lagrangian** for the $T + 1$ period problem.
- However, a **better way** would be to somehow make use of the T period solution that we found, $V^T(W_1)$ to create $V^{T+1}(W_0)$
- Dynamic Programming (DP) provides means for doing this.
- **DP essentially converts a general T period problem into a 2 period one.**

Cake-Eating Example - Dynamic Programming Approach

- DP breaks down the optimal path into two parts, what is **optimal today** and the **optimal continuation path**.
- Given W_0 , the optimization problem can be written as:

$$V^{T+1}(W_0) = \max_{\{c_t, W_{t+1}\}_{t=0}^T} \left[\sum_{t=0}^T \beta^t u(c_t) \right] \quad (8)$$

$$= \max_{\{c_t, W_{t+1}\}_{t=0}^T} \left[u(c_0) + \sum_{t=1}^T \beta^t u(c_t) \right]$$

$$= \max_{\{c_t, W_{t+1}\}_{t=0}^T} \left[u(c_0) + \beta \sum_{t=1}^T \beta^{t-1} u(c_t) \right]$$

$$= \max_{c_0, W_1} \left[u(c_0) + \beta \max_{\{c_t, W_{t+1}\}_{t=1}^T} \left[\sum_{t=1}^T \beta^{t-1} u(c_t) \right] \right]$$

$$V^{T+1}(W_0) = \max_{c_0, W_1} [u(c_0) + \beta V^T(W_1)] \quad (9)$$

Cake-Eating Example - Dynamic Programming Approach

- Subject to the constraint

$$W_1 = W_0 - c_0$$

- Note V^T here denotes value function for the **T -periods-left model** not value function at time T !!! Best to think of this as V_t^T for some time t .
- In terms of time t notation, the general **Bellman equation** is:

$$V_t(W_t) = \max_{c_t, W_{t+1}} \left\{ u(c_t) + \beta V_{t+1}(W_{t+1}) \right\}$$

where $t = 0, 1, \dots, T$.

- This is a **functional equation** - the unknown is now a function V .

Cake-Eating Example - Dynamic Programming Approach

- So instead of choosing the entire path of c_t , we are just choosing c_0 .
- The rest of the path is optimally determined by the value function, $V^T(W_1)$.
- Once c_0 and hence W_1 is determined, the value function summarizes the rest of the problem
- This is the **principle of optimality** due to **Richard Bellman**: we can represent the full dynamic problem as a sequence of recursive 2 period problems:
- Optimal Today + Optimal Continuation Path.

Cake-Eating Example - Dynamic Programming Approach

- The **Bellman equation** for the cake eating problem is then written as

$$V_t(W_0) = \max_{c_t, W_{t+1}} [u(c_t) + \beta V_{t+1}(W_1)]$$

where $t = 0, 1, \dots, T$. Here V_t is the value function at any time t and V_{t+1} is the value function for the next period $t + 1$.

- The **solution** to this problem is given by the **decision rules (functions)** for consumption and next period cake size: $c_t(W_t)$ and $W_{t+1}(W_t)$.
- To obtain these decision rules, we need to find the **unknown value function** $V_t(W_t)$.
- Since this is a finite horizon problem, we can achieve this task easily. Start with the last period T where $V_{T+1} = 0$ and work backwards to obtain all the other value functions and decision rules.

Cake-Eating Example - Dynamic Programming Approach

Substituting for W_{t+1} from the constraint, we can write eq. (7) as:

$$V_t(W_t) = \max_{c_t} [u(c_t) + \beta V_{t+1}(W_t - c_t)]$$

The **first order condition** of this value function problem [EC] is given by:

$$u'(c_t) = \beta V'_{t+1}(W_t - c_t)$$

Denote the solution to the problem, optimal consumption by $c_t^* = h_t(W_t)$.

Then the value function is

$$V_t(W_t) = [u(h_t(W_t)) + \beta V_{t+1}(W_t - h_t(W_t))]$$

Taking the derivative w.r.t W_t , we get the **Envelope condition**¹

$$V'_t(W_t) = [u'(h_t(W_t))h'_t(W_t) + \beta V'_{t+1}(\cdot)[1 - h'_t(W_t)]] \quad (10)$$

$$= u'(c_t) \quad (11)$$

¹borrow the FOC for the second term sub

Cake-Eating Example - Dynamic Programming Approach

Taking one period forward, with stationarity of the value function:

$$V'_{t+1}(W_{t+1}) = u'(c_{t+1})$$

The FOC along with the above envelope condition together imply the Euler equation,

$$u'(c_t) = \beta u'(c_{t+1}) \quad \text{for } t = 0, 1, 2, \dots, T - 1$$

So the **solution is the same** whether we use the **recursive DP approach** or the **sequential Lagrangian** method.

Cake-Eating Example - Infinite Horizon

- **Suppose we allow the horizon to go to infinity.**
- As before, one can consider solving the infinite horizon sequence problem given by:

$$\max_{\{c_t\}_0^\infty, \{W_{t+1}\}_0^\infty} \sum_{t=0}^{\infty} \beta^t u(c_t)$$

along with the transition equation of

$$W_{t+1} = W_t - c_t$$

for $t = 0, 1, 2, \dots \infty$ and some given $W_0 > 0$.

Infinite Horizon

Infinite Horizon - Dynamic Programming

- Since the time horizon is infinite, the future from today and the future from tomorrow is of the same length (which is infinity).
- Therefore, the value function for the problem, which is the sum of discounted utility that the agent gains optimally, is not a function of the time period, but only of the cake size.
- The value function for the infinite horizon case is

$$V(W_t) = \max_{\{c_t\}_0^\infty, \{W_t\}_1^\infty} \sum_{t=0}^{\infty} \beta^t u(c_t)$$

Infinite Horizon - Dynamic Programming

We can form the **Bellman equation** by breaking down this infinite sequence into a **recursive** two-period problem:

$$V(W_t) = \max_{\{c_t\}_0^\infty, \{W_t\}_1^\infty} \sum_{t=0}^{\infty} \beta^t u(c_t) \quad (12)$$

$$= \max_{\{c_t\}_0^\infty, \{W_t\}_1^\infty} \left[u(c_0) + \sum_{t=1}^{\infty} \beta^t u(c_t) \right] \quad (13)$$

$$= \max_{c_0, W_1} \left[u(c_0) + \max_{\{c_t, W_{t+1}\}_1^\infty} \sum_{t=1}^{\infty} \beta^t u(c_t) \right] \quad (14)$$

$$= \max_{c_0, W_1} \left[u(c_0) + \beta \max_{\{c_t, W_{t+1}\}_1^\infty} \sum_{t=0}^{\infty} \beta^t u(c_{t+1}) \right] \quad (15)$$

$$V(W_t) = \max_{c_0, W_1} [u(c_0) + \beta V(W_{t+1})] \quad (16)$$

- So the **infinite horizon dynamic programming problem** is

$$V(W) = \max_{c, W'} \left\{ u(c) + \beta V(W') \right\} \quad \text{for all } W \quad (17)$$

$$s.t. \quad W' = W - c \quad (18)$$

- Variables with **prime** denote **future values**².
- $V(W)$ is the **value** of the infinite horizon cake eating problem or the **maximal utility from this consumption**.
- $W' = W - c$ is the **state transition equation** or equivalently the evolution of cake size.

²not to be confused with derivatives, that is W denotes W_t and W' denotes W_{t+1}

- In general, we use **primes** to denote **future values** when we are looking for a stationary solution to an infinite horizon problem.
- The value function here is **stationary**, that is:

$$V_t(W) = V_{t+k}(W) = V(W) \quad \text{for any } k > 0$$

- **Stationarity** means time-invariant, that is the value function or policy functions are optimal and do not change with time.

Infinite Horizon - Remarks

- Remember these functions denote **a path or a rule**, so **stationarity** here means that this **path is constant** (not the actual variable).
- The two policy functions maps the state variables to controls.
- In this problem, the two policy functions are:

$$W'(W) \text{ and } c(W)$$

next period cake size and consumption.

Infinite Horizon - State and Control Variables

- What are the state and control (choice) variables?
- The **state variable** is the size of the cake (W) that is given at the start of any period.
- The cake size completely summarizes all information from the past that is needed for the forward looking optimization problem.
- The **control variable** is the variable that is being **chosen**. In this case, it is the level of consumption in the current period, c and next period cake size W' .
- The **transition (or the constraint)** describes the dependence of the state tomorrow on the state today and the control today:

$$W' = W - c$$

Infinite Horizon - State and Control Variables

- Alternatively, we can write the DP, in (10), as:

$$V(W) = \max_{W'} \left\{ u(W - W') + \beta V(W') \right\}$$

where we have substituted the constraint so that we have to choose only tomorrow's cake size.

- **Either specification will yield the same result.** Fewer choice variables are easier to work with.
- This expression is a **functional equation** and is often called a **Bellman equation after Richard Bellman**, the originator of dynamic programming.
- Note that the **unknown in the Bellman equation is the value function itself**: the idea is to find a function $V(W)$ that satisfies this condition for all W .

Items for Review

- Sequential Lagrangian
- Shadow Price
- Consumption/Saving with no production, depreciation
- Sequential solution with Euler Equation
- Shooting Algorithm
- Finite Horizon
- Recursive Approach
- Bellman Equation
- Continuation Value
- Infinite Horizon
- State and Choice/Control Variables